



Models@Run.Time Workshop 2015

# USING REFERENCE ATTRIBUTE GRAMMAR-CONTROLLED REWRITING FOR ENERGY AUTO-TUNING

Christoff Bürger    Department of Computer Science, Lund University, Sweden

**Johannes Mey**    Software Technology Group, TU Dresden, Germany

René Schöne    Software Technology Group, TU Dresden, Germany

Sven Karol    Chair for Compiler Construction, TU Dresden, Germany

Daniel Langer    Software Technology Group, TU Dresden, Germany

Ottawa, September 29, 2015

# Presentation Overview

Our **new idea**: Use Reference Attribute Grammars and rewriting for runtime models.

We use

- a **Reference Attribute Grammar (RAG)**
- to create and modify a **runtime model**
- for batch process execution on a compute cluster and
- use **attributes** and **RAG-controlled rewrites** to schedule the system's tasks
- in an **energy-optimized** way.

# Outline

## Case Study

## Solution Background (*RACR*)

## Our Solution

## Evaluation and Outlook

# A Case Study: Scheduling Batch Processes

## of Wikipedia Indexing Tasks

Very simple case study to show use of RAGs for runtime model


**Task:** indexing of text chunks (taken from Wikipedia)

- predictable processing time (proportional to chunk size)
- requests arrive interactively (occur randomly)
- requests have deadline


**Energy Optimization:** Minimize energy consumption of the indexing system

- System is network of (embedded) computers
- Computers (and connecting switches) can be turned off to save energy


# A Case Study: Scheduling Batch Processes of Wikipedia Indexing Tasks



# A Case Study: Scheduling Batch Processes of Wikipedia Indexing Tasks



# A Case Study: Scheduling Batch Processes of Wikipedia Indexing Tasks



# Solution Background

## RACR - Reference Attribute Grammar-Controlled Rewriting

RACR is ...

- a Reference Attribute Grammar (RAG) system
  - declarative semantics
  - lazy, incremental evaluation
- for RAG-controlled rewriting
  - advanced AST manipulation



# Solution Background

## RACR - Reference Attribute Grammar-Controlled Rewriting

RAG-controlled rewriting = RAGs + graph rewriting

- reference attribute grammar for declarative **analyses**
  - reference attributes induce semantic overlay graph on top of *abstract syntax tree (AST)*
  - enables deduction *and* analyses of graph structure
  - deduced, memoized *abstract syntax graph (ASG)*
- graph rewriting for ASG **transformations**
  - left hand: ASG pattern (ASTs connected via reference attributes)
  - right hand: manipulations on matched underlying AST
  - ASG changes with AST (updated by RAG)
- seamless combination:
  - use analyses to deduce rewrites
  - rewrites automatically update analyses
  - incremental

# Solution Background

## RACR - Reference Attribute Grammar-Controlled Rewriting

### *The Implementation: RACR*

- reference implementation of RAG-controlled rewriting in *Scheme* R6RS<sup>1</sup>

*RACR* contains API for:

- ASG schema definition (AST schema + attribution)
- ASG querying (AST + attributes)
- rewriting:
  - imperative **and/or** RAG-controlled **and/or** fixpoint
  - primitive **and/or** pattern-based
  - ... **in any combination!**

<https://github.com/christoff-buerger/racr>

# Solution Background

## RACR - Reference Attribute Grammar-Controlled Rewriting

### *The Implementation: RACR*

- reference implementation of RAG-controlled rewriting in *Scheme* R6RS<sup>1</sup>

*RACR* contains API for:

- ASG schema definition (AST schema + attribution)
- ASG querying (AST + attributes)
- rewriting:
  - imperative **and/or** RAG-controlled **and/or** fixpoint
  - primitive **and/or** pattern-based
  - ... **in any combination!**

<https://github.com/christoff-buerger/racr>

<sup>1</sup>Don't panic! C bindings and .NET bindings are available.

# Our Solution

## The Grammar

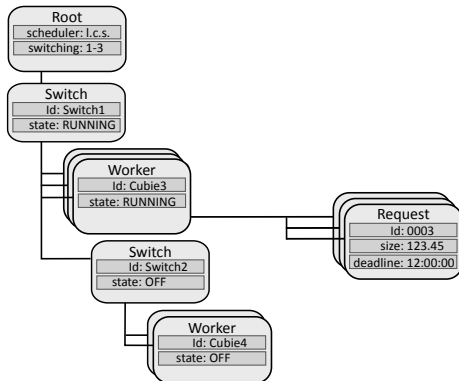
Grammar is encoded in Scheme symbols

- production rule: left side  $\rightarrow$  right side
- upper case: nonterminals
- lower case: terminals
- repetition (\*), inheritance (:)

```
(ast-rule 'Root->scheduler-backupworkers-CompositeWorker)
(ast-rule 'AbstractWorker->id-state-timestamp)
(ast-rule 'CompositeWorker:AbstractWorker->AbstractWorker*)
(ast-rule 'Switch:CompositeWorker->)
(ast-rule 'Worker:AbstractWorker->devicetype-Request*<Queue)
(ast-rule 'Request->id-size-deadline-dispatchtime)
```

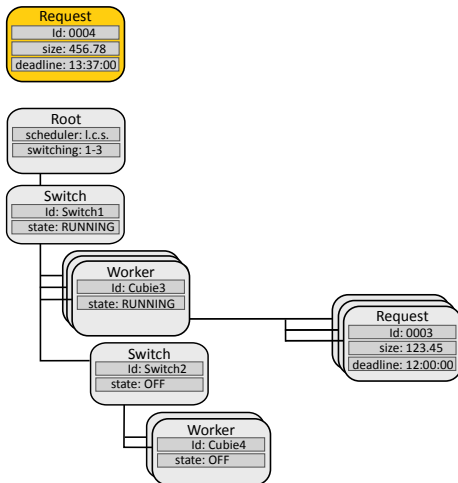
# Our Solution

## Example AST



# Our Solution

## Scheduling a Request



# Our Solution

## Scheduling by Rewriting

Scheduling a new task: rewriting the AST

- insert a new Request node at the right position

```
(rewrite-insert  
  (ast-child 'Queue worker) ;list-node to insert into  
  index ;position of insertion  
  (create-ast spec 'Request (list id size deadline #f)))
```

# Our Solution

## Attribute-controlled Scheduling

Where to put the new Request?

- evaluate attribute **schedule** to find insertion position
- result is worker and position in worker's queue
- Attribute depends on terminal **scheduler**  
→ scheduler can be exchanged at runtime!

```
(ag-rule schedule
  (Root (lambda (n time work-id load-size deadline)
    (att-value (ast-child 'scheduler n) n time work-id
      load-size deadline))))
```

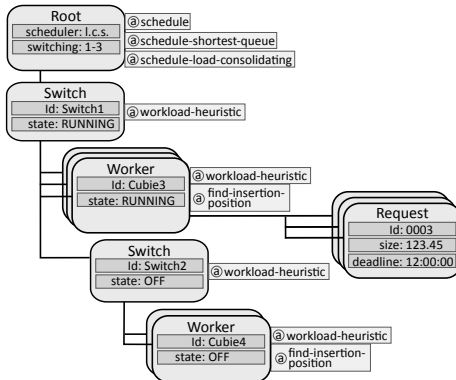


# Our Solution

Two schedulers implemented:

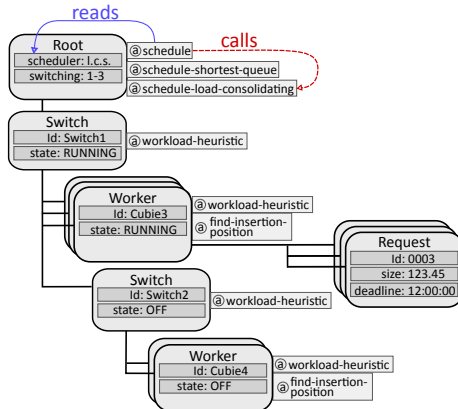
**schedule-shortest-queue** simple scheduler inserting in shortest queue of any worker

**schedule-load-consolidating** inserts request in fullest queue while ensuring deadline is kept



# Our Solution

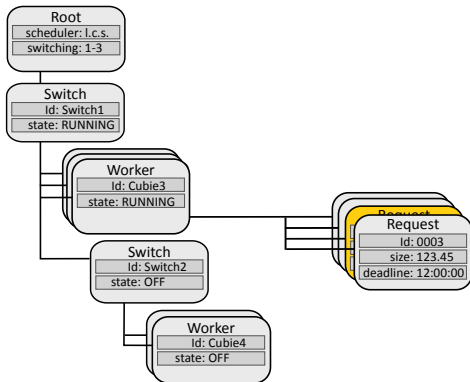
## Attribute evaluation



# Our Solution

## Scheduling a Request

### The resulting AST



# Our Solution

## Saving Energy

Required workers are computed with attributes

- adaptation strategy regularly computes *how many* and *which* workers to switch on or off
- interactive system requires backup workers
- amount of backup workers and adaptation parameters described in AST

Saving energy by switching off workers:

- Try to minimize amount of idle workers
- use adaptation strategy
- use load-consolidating scheduler to minimize required workers

# Evaluation

Test setup for measuring energy consumption

- graphical interface to show system state and consumed power
- Scenario generator to run controlled workloads with different settings


## Evaluation

**Requests**

Shortest-queue scheduler, workers always on



## Evaluation



## Shortest-queue scheduler, workers always on



## Evaluation


[illegible]

## Energy-aware shortest-queue scheduler





# Evaluation



## Energy-aware shortest-queue scheduler




## Evaluation

[illegible]

## Energy-aware load-consolidating scheduler



# Evaluation



## Energy-aware load-consolidating scheduler



# Evaluation

## Properties of the Solution

- **scalable**: incremental evaluation ensures only necessary attributes are re-evaluated after system change
- **adaptive**: ASG structure can be modified at runtime, schedulers and parameters can be switched
- **fault-tolerant**: system can handle device failures

# Evaluation

## Results

- Energy-aware shortest queue scheduler saves 13.1% compared to regular shortest-queue scheduler
- Energy-aware load-consolidating scheduler saves 17.5% compared to regular shortest-queue scheduler
- load-consolidating scheduler increases amount of request that can be scheduled

# Outlook

## Next steps

- heterogeneous architecture
  - more interesting network structure
  - simulate large systems
- 
- more case studies for RACR for runtime models

# Conclusion

## Benefits of RAG-Controlled Rewriting

